



Computational Tools for Macroeconometrics

Assignment 3

This assignment covers two topics: Monte Carlo simulations and Bootstrap, with a focus on time series models. As usual, I will discuss the basic ideas behind the concepts and then present the assignment.

Monte Carlo experiments

Monte Carlo experiments is a handy tool to assess the quality of the asymptotic approximation of econometric estimators.

Consider the following linear model

$$y_t = \beta_0^o + \beta_1^o x_t + u_t, t = 1, \dots, T,$$

where y_t , x_t , and u_t are random variables and β_0^o and β_1^o are parameters to be estimated. In matrix form, the model can be written as

$$Y_t = \mathbf{X}_t \beta^o + U_t,$$

where

$$\underline{Y}_{(T \times 1)} = \begin{pmatrix} Y_1 \\ \vdots \\ Y_T \end{pmatrix}, \underline{X}_{T \times 2} = \begin{pmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_T \end{pmatrix}, \underline{U}_{(T \times 1)} = \begin{pmatrix} u_1 \\ \vdots \\ u_T \end{pmatrix}, \underline{\beta}_{(2 \times 1)}^o = \begin{pmatrix} \beta_0^o \\ \beta_1^o \end{pmatrix}$$

We will assume that 1. (u_t, x_t) are i.i.d. over t ; 2. $E(u_t | x_t) = 0, t = 1, \dots, T$; 3. $E(|x_t|^4) < \infty$ and $E(|u_t|^4) < \infty$; 4. $E(u^t | x_t) = \sigma^2 > 0$.

Under these assumptions, the OLS estimator of β^o ,

$$\hat{\beta}^{ols} = (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'Y$$

is unbiased, consistent, and asymptotically normal:

$$E(\hat{\beta}^{ols}) = \beta^o, \quad \hat{\beta}^{ols} \xrightarrow{p} \beta^o, \quad \left(\hat{\beta}^{ols} - \beta^o\right) \xrightarrow{d} N\left[0, \sigma^2 E(\mathbf{X}'\mathbf{X})^{-1}\right].$$

These results are derived theoretically. For unbiasedness, the law of iterated expectations and the iid of (u_t, x_t) gives

$$E\left(E\left(\hat{\beta}^{ols}|\mathbf{X}\right)\right) = \beta^o + E\left((\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'E(U|\mathbf{X})\right) = \beta^o.$$

Under the iid assumption and the restrictions on the fourth moments of u_t and x_t , we have

$$\frac{1}{T}\mathbf{X}'U = \frac{1}{n}\sum_{t=1}^T \begin{pmatrix} u_t \\ x_t u_t \end{pmatrix} \xrightarrow{p} \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \text{ and } \frac{1}{T}\mathbf{X}'\mathbf{X} = \frac{1}{T}\sum_{t=1}^T \begin{pmatrix} 1 & x_t \\ x_t & x_t^2 \end{pmatrix} \xrightarrow{p} \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}.$$

These convergences in probability deliver consistency of the OLS estimator

$$\hat{\beta}^{ols} = \beta^o + \left(\frac{1}{T}\mathbf{X}'\mathbf{X}\right)^{-1} \frac{1}{T}\mathbf{X}'U \xrightarrow{p} \beta^o + \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}^{-1} \begin{pmatrix} 0 \\ 0 \end{pmatrix} = \beta^o.$$

Finally, iid, the moments' restrictions, and the homoskedasticity assumptions, give

$$\sqrt{T}\left(\hat{\beta}^{ols} - \beta^o\right) \xrightarrow{d} N\left[0, \sigma^2 \begin{pmatrix} 1 & E(x_t) \\ E(x_t) & E(x_t^2) \end{pmatrix}^{-1}\right].$$

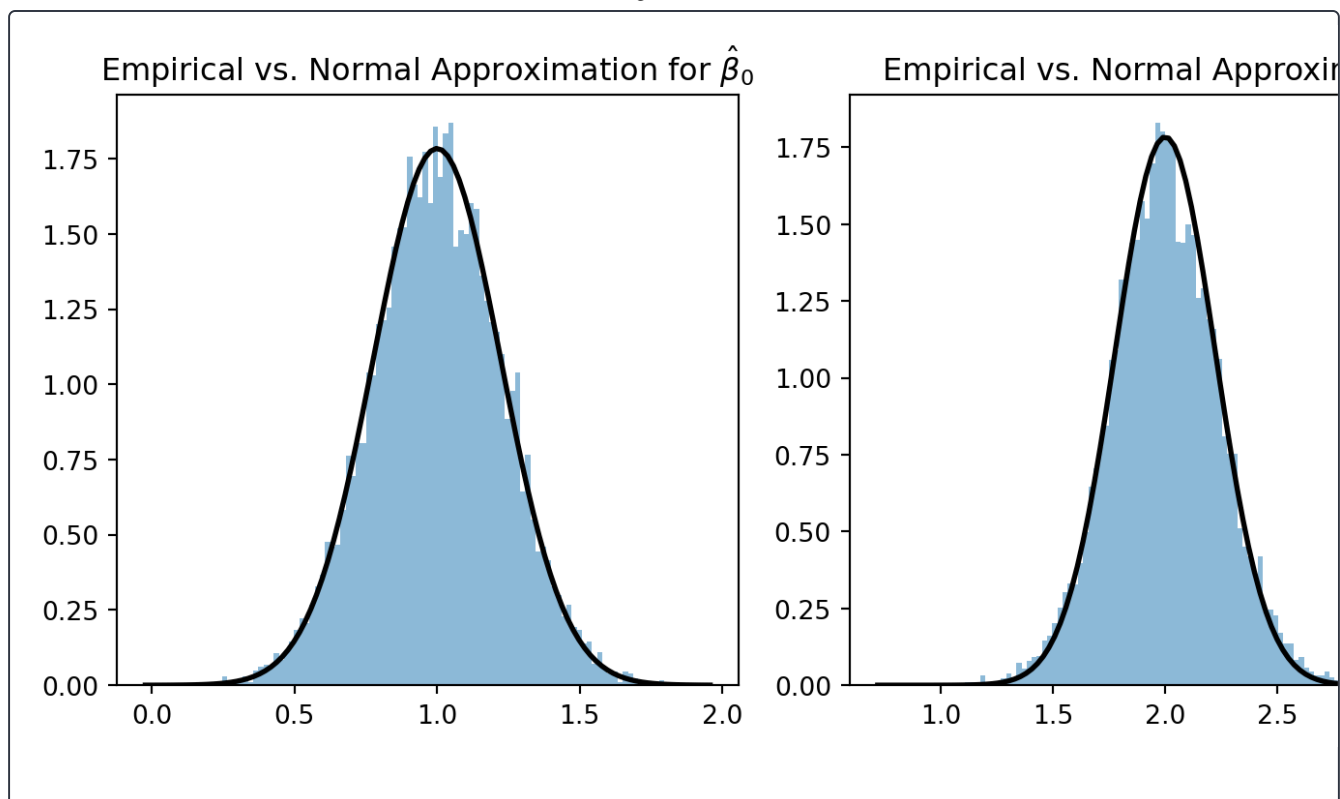
Consistency and asymptotic results, and they hold as $T \rightarrow \infty$. Given a sample of size T , say $T = 50$, how close is the distribution of $\hat{\beta}^o$ the normal one postulated by the usual asymptotic theory? How close is $\hat{\beta}^o$ to β^o ? Unfortunately, theory only tells us that the larger the sample size, the better the normal approximation. Similarly, it tells us that $\hat{\beta}^o - \beta^o$ is smaller the larger the sample size.

Monte Carlo simulations can help us determine how good these approximations are for a given size of T . The idea is simple: we simulate a sample of (u_t, x_t) , $t = 1, \dots, T$ from which we generate y_t (using a arbitrary value for β_0^o and β_1^o .) We estimate the parameters using the simulated data. We repeat this operation many times, saving the estimated parameters. The saved parameters give the empirical distribution of the OLS estimator and help verify how closely the theory matches the empirical distribution.

The following code does a Monte Carlo for the linear model above:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.stats import norm
4
5 # Set parameters
6 T = 20
```

```
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15
16 # Run simulations
17 for i in range(num_simulations):
18     x = np.random.normal(0, 1, T)
19     u = np.random.normal(0, sigma, T)
20     y = beta_0_true + beta_1_true * x + u
21
22     # OLS estimation
23     X = np.vstack([np.ones(T), x]).T
24     beta_hat = np.linalg.inv(X.T @ X) @ X.T @ y
25     beta_0_estimates[i] = beta_hat[0]
26     beta_1_estimates[i] = beta_hat[1]
27
28 # Plotting the results
29 fig, ax = plt.subplots(1, 2, figsize=(8, 4))
30
31 # Distribution of beta_0
32 ax[0].hist(beta_0_estimates, bins=100, alpha=0.5, density=True)
33 xmin, xmax = ax[0].get_xlim()
34 x = np.linspace(xmin, xmax, 100)
35 p = norm.pdf(x, beta_0_true, 1/np.sqrt(T))
36 ax[0].plot(x, p, 'k', linewidth=2)
37 ax[0].set_title(f'Empirical vs. Normal Approximation for  $\hat{\beta}_0$ ')
38
39 # Distribution of beta_1
40 ax[1].hist(beta_1_estimates, bins=100, alpha=0.5, density=True)
41 xmin, xmax = ax[1].get_xlim()
42 x = np.linspace(xmin, xmax, 100)
43 p = norm.pdf(x, beta_1_true, 1/np.sqrt(T))
44 ax[1].plot(x, p, 'k', linewidth=2)
45 ax[1].set_title(f'Empirical vs. Normal Approximation')
46
47 plt.tight_layout()
48 plt.show()
```



The approximation is excellent, even if $T = 20$. This might sound surprising since the normal distribution of $\hat{\beta}^{ols}$ is normal when T is large. We get such a close agreement between the empirical distribution of $\hat{\beta}^{ols}$ and the theoretical one because we are simulating u_t the data from a normal distribution independently from x_t . When $u_t|x_t \sim N(0, \sigma^2)$, the distribution of the OLS estimator is precisely normal, even when $T = 3$. Of course, if T is small, the estimator's variance will be larger, but a normal distribution will approximate the OLS estimators' distribution.

Instead of squinting at the histograms and the density implied by the CLT, we can modify the code to calculate the confidence interval at each simulation and see how many times the "true" values of the parameters fall into all the intervals generated. If the approximation is good, a 95% confidence interval should contain the "true" parameters 95% of the time.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.stats import norm
4
5  # Set parameters
6  T = 50
7  beta_0_true = 1
8  beta_1_true = 2
9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)

```

```

15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = np.random.normal(0,1,T)
21     u = np.random.normal(0,sigma,T)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check weather beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.9
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.9
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

```

Empirical 95% CI for beta_0: 0.9457
Empirical 95% CI for beta_1: 0.9428

```

The empirical coverage of the confidence intervals is only in the neighborhood of 95%. The reason is that we are estimating the variance of the coefficient, and in this case, the distribution of $\hat{\beta}_1/se(\hat{\beta}_1)$ has a t -student distribution with $T - 1$ degrees of freedom. If you run the code above for $T = 50$, the confidence interval coverage will be closer to 95%.

Let us see what happens if we simulate x_t and u_t from a chi-squared distribution with 5 of freedom centered in a way to have mean zero and unit variance.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.stats import norm
4
5 # Set parameters
6 T = 50
7 beta_0_true = 1
8 beta_1_true = 2
9 sigma = 1

```

```

10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = (np.random.chisquare(4,T) - 4)/np.sqrt(2*4)
21     u = (np.random.chisquare(4,T) - 4)/np.sqrt(2*4)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check weather beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.9
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.9
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

```

Empirical 95% CI for beta_0: 0.9397
Empirical 95% CI for beta_1: 0.9485

```

Even a $T = 50$, there is some difference between the empirical and theoretical coverage.

We can use the simulations to see what happens when the conditions on the moments of u_t and x_t are unsatisfied.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.stats import norm
4
5 # Set parameters
6 T = 150
7 beta_0_true = 1
8 beta_1_true = 2

```

```

9  sigma = 1
10 num_simulations = 10000
11
12 # Arrays to store the estimates from each simulation
13 beta_0_estimates = np.zeros(num_simulations)
14 beta_1_estimates = np.zeros(num_simulations)
15 beta_0_in = np.zeros(num_simulations)
16 beta_1_in = np.zeros(num_simulations)
17
18 # Run simulations
19 for i in range(num_simulations):
20     x = np.random.standard_cauchy(T)
21     u = np.random.standard_cauchy(T)
22     y = beta_0_true + beta_1_true * x + u
23     # OLS estimation
24     X = np.vstack([np.ones(T), x]).T
25     XXinv = np.linalg.inv(X.T @ X)
26     beta_hat = XXinv @ X.T @ y
27     beta_0_estimates[i] = beta_hat[0]
28     beta_1_estimates[i] = beta_hat[1]
29     u_hat = y - beta_hat[0] - beta_hat[1] * x
30     sigma2_hat = np.dot(u_hat, u_hat)/(T-2)
31     variance_hat = sigma2_hat*XXinv
32     se_0 = np.sqrt(variance_hat[0,0])
33     se_1 = np.sqrt(variance_hat[1,1])
34     ## Check weather beta_0 in CI 95%
35     beta_0_in[i] = beta_hat[0] - 1.965*se_0 < beta_0_true < beta_hat[0] + 1.9
36     beta_1_in[i] = beta_hat[1] - 1.965*se_1 < beta_1_true < beta_hat[1] + 1.9
37
38 # Output the results
39 print(f"Empirical 95% CI for beta_0: {np.mean(beta_0_in)}")
40 print(f"Empirical 95% CI for beta_1: {np.mean(beta_1_in)}")

```

Empirical 95% CI for beta_0: 0.9764

Empirical 95% CI for beta_1: 0.9541

Bootstrap

The *bootstrap* is a statistical tool for estimating the distribution of a statistic based on random sampling with replacement. It allows for assessing a statistic's variability and is beneficial in settings where the theoretical distribution of the statistic is unknown or difficult to derive.

In the context of a linear regression model, we can use the bootstrap to perform inference on the regression coefficients. Despite its theoretical intricacies, the

bootstrap is conceptually simple and describing its application can be done quite easily.

Input: Dataset (X, y) with n observations
 Fit original model: Estimate coefficients β from original dataset
 Store these coefficients
 Bootstrap procedure: For $b = 1$ to \$(number of bootstrap samples):
 Create "bootstrap" sample (X^*, y^*) by randomly sampling n observations with replacement from original dataset
 Fit linear model to bootstrap sample
 Estimate and store bootstrap coefficients β^*
 Calculate statistics: Compute variance of bootstrap coefficient estimates
 Optionally calculate confidence intervals using percentiles of bootstrap distribution
 Output: Original coefficients, bootstrap variance estimates, confidence intervals

```

1  import numpy as np
2
3  # Set random seed for reproducibility
4  np.random.seed(0)
5
6  # Generate some sample data
7  T = 100 # Number of observations
8  x = np.random.normal(0, 1, T)
9  u = np.random.normal(0, 1, T)
10 beta_0_true = 1
11 beta_1_true = 2
12
13 # Simulate response variable y
14 y = beta_0_true + beta_1_true * x + u
15
16 # Function to fit linear model
17 def fit_linear_model(x, y):
18     X = np.vstack([np.ones(len(x)), x]).T
19     beta_hat = np.linalg.inv(X.T @ X) @ (X.T @ y)
20     return beta_hat
21
22 # Initial fit
23 initial_beta = fit_linear_model(x, y)
24
25 # Number of bootstrap samples
26 B = 1000
27 bootstrap_estimates = np.zeros((B, 2))
28
29 # Perform bootstrap resampling
30 for i in range(B):
31     indices = np.random.choice(range(T), size=T, replace=True)
32     x_resampled = x[indices]
33     y_resampled = y[indices]
34     bootstrap_estimates[i] = fit_linear_model(x_resampled, y_resampled)
35
36 # Compute standard errors

```



```

37 standard_errors = bootstrap_estimates.std(axis=0)
38
39 print("Bootstrap Standard Errors:")
40 print("SE(beta_0):", standard_errors[0])
41 print("SE(beta_1):", standard_errors[1])
42
43 print("LM Standard Errors:")
44 import statsmodels.api as sm
45 X = sm.add_constant(x)
46 model = sm.OLS(y, X)
47 results = model.fit()
48
49 # Standard errors from statsmodels
50 statsmodels_se = results.bse
51 print("Standard Errors from statsmodels OLS:")
52 print("SE(beta_0):", statsmodels_se[0])
53 print("SE(beta_1):", statsmodels_se[1])

```

```

Bootstrap Standard Errors:
SE(beta_0): 0.10061319113712466
SE(beta_1): 0.09469134082416591
LM Standard Errors:
Standard Errors from statsmodels OLS:
SE(beta_0): 0.10404543947505016
SE(beta_1): 0.10305046686003079

```

Linear model with dependent data

So far, we have assumed that (u_t, x_t) are iid. With macroeconomic data, it is often the case that (u_t, x_t) are correlated.

Consider the following model:

$$y_t = \beta_0 + \beta_1 x_t + u_t,$$

with

$$x_t = \phi_x x_{t-1} + \eta_t, \quad |\phi_x| < 1$$

and

$$u_t = \phi_u u_{t-1} + \varepsilon_t, \quad |\phi_u| < 1.$$

If η_t and ε_t are independent, we will have that $E(u_t | x_t) = 0$ and so we can consistently estimate β_1 by OLS. The asymptotic distribution of the OLS estimator will be normal, but the variance of this distribution is difficult to estimate. If we use the standard errors

that do not consider the correlation over time of the variables, the inference will be off. We perform a simple Monte Carlo to see the impact of serial correlation in x_t and u_t .

```

1  def simulate_ar1(n, phi, sigma):
2      """
3      Simulate an AR(1) process.
4
5      Parameters:
6      n (int): Number of observations.
7      phi (float): Coefficient of AR(1) process.
8      sigma (float): Standard deviation of the innovation term.
9
10     Returns:
11     np.array: Simulated AR(1) error terms.
12     """
13     errors = np.zeros(n)
14     eta = np.random.normal(0, sigma, n) # white noise
15     for t in range(1, n):
16         errors[t] = phi * errors[t - 1] + eta[t]
17     return errors
18
19 def simulate_regression_with_ar1_errors(n, beta0, beta1, phi_x, phi_u, sig
20     """
21     Simulate a regression model with AR(1) error terms.
22     Parameters:
23     n (int): Number of observations.
24     beta0 (float): Intercept of the regression model.
25     beta1 (float): Slope of the regression model.
26     phi (float): Coefficient of the AR(1) process in the error term.
27     sigma (float): Standard deviation of the innovation term in the AR(1)
28     Returns:
29     tuple: x (independent variable), y (dependent variable), errors (AR(1) p
30     """
31     x = simulate_ar1(n, phi_x, sigma)
32     u = simulate_ar1(n, phi_u, sigma)
33     y = beta0 + beta1 * x + u
34     return x, y, u
35
36 T = 500                # Number of observations
37 beta0 = 1.             # Intercept
38 beta1 = 2              # Slope
39 phi_x = 0.7            # AR(1) coefficient for x
40 phi_u = 0.7            # AR(1) coefficient for the errors
41 sigma = 1              # Standard deviation of the white noise
42
43 # Simulating the model
44
45 ## Do monte carlo
46 t_stats_hc = []
47 t_stats_hac = []

```

```

48
49 for i in range(1000):
50     x, y, errors = simulate_regression_with_ar1_errors(T, beta0, beta1, phi_
51     X = sm.add_constant(x)
52     model = sm.OLS(y, X).fit(cov_type='HC1')
53     t_stats_hc.append(model.t_test('x1=2').tvalue)
54     ## Use HAC: takes into account serial correlation
55     model2 = sm.OLS(y, X).fit(cov_type='HAC', cov_kws={'maxlags': np.floor(
56     t_stats_hac.append(model2.t_test('x1=2').tvalue)
57
58 ## Check we reject the null hypothesis at alpha=0.05 about 5% of the time
59
60 print(f"Empirical size test beta_1=2 using White SE: {np.mean(np.abs(np.ar
61 print(f"Empirical size test beta_1=2 using HAC SE: {np.mean(np.abs(np.arra

```

```

Empirical size test beta_1=2 using White SE: 0.255
Empirical size test beta_1=2 using HAC SE: 0.075

```

We can use the bootstrap to obtain the standard errors of $\hat{\beta}_1$. Unfortunately, with time-dependent data, the bootstrap needs to be modified. Instead of sampling with replacement observations from y_t and x_t , we will sample blocks of length ℓ with replacement. By resampling blocks, we ensure that the correlation in the data is preserved.

The moving block bootstrap (MBB) adapts the traditional bootstrap method to handle data where observations are dependent, such as in time series analysis. This approach involves resampling consecutive observation blocks to preserve the data's internal structure and dependence.

Steps for Moving Block Bootstrap (MBB):

1. **** Choose Block Length ****: Determine the length of the blocks, ℓ , that will be resampled. This length should be chosen based on the data's correlation structure: it should be large enough to capture the dependence within the data.
2. **Generate Blocks**: From the original dataset of size T , generate new datasets by sampling blocks of length ℓ and concatenate them until reaching the size T . Blocks should be sampled with replacement.
3. **Resample Data within Blocks**: Sample y_t and x_t using the sampled blocks. This is the bootstrapped dataset.
4. **Refit the Model**: Fit the regression model to each bootstrapped dataset and collect the estimated parameters for each dataset.

5. Calculate Statistics: Calculate the standard deviation of the bootstrap estimate.

These are the standard errors.

```

1 import numpy as np
2 import statsmodels.api as sm
3
4 def moving_block_bootstrap(x, y, block_length, num_bootstrap):
5     T = len(y) # Total number of observations
6     num_blocks = T // block_length + (1 if T % block_length else 0)
7
8     # Fit the original model
9     X = sm.add_constant(x)
10    original_model = sm.OLS(y, X)
11    original_results = original_model.fit()
12
13    bootstrap_estimates = np.zeros((num_bootstrap, 2)) # Storing estimates
14
15    # Perform the bootstrap
16    for i in range(num_bootstrap):
17        # Create bootstrap sample
18        bootstrap_indices = np.random.choice(np.arange(num_blocks) * block_length,
19                                             num_bootstrap)
20        bootstrap_sample_indices = np.hstack([np.arange(index, min(index + block_length, T))
21                                             for index in bootstrap_indices])
22        # Ensure the sample indices are within the range of the data
23        bootstrap_sample_indices = bootstrap_sample_indices[:T]
24
25        x_bootstrap = x[bootstrap_sample_indices]
26        y_bootstrap = y[bootstrap_sample_indices]
27
28        # Refit the model on bootstrap sample
29        X_bootstrap = sm.add_constant(x_bootstrap)
30        bootstrap_model = sm.OLS(y_bootstrap, X_bootstrap)
31        bootstrap_results = bootstrap_model.fit()
32
33        # Store the estimates
34        bootstrap_estimates[i, :] = bootstrap_results.params
35
36    return bootstrap_estimates
37
38 # Run moving block bootstrap
39 block_length = 12
40 num_bootstrap = 1000
41 x, y, errors = simulate_regression_with_ar1_errors(200, beta0, beta1, phi)
42 bootstrap_results = moving_block_bootstrap(x, y, block_length, num_bootstrap)
43
44 # Calculate and print standard errors
45 bootstrap_standard_errors = bootstrap_results.std(axis=0)
46 print("Bootstrap Standard Errors:")
47 print("SE(beta_0):", bootstrap_standard_errors[0])
48 print("SE(beta_1):", bootstrap_standard_errors[1])

```

```
Bootstrap Standard Errors:  
SE(beta_0): 0.026787210831234053  
SE(beta_1): 0.05979213797333231
```

Assignment 3

Task

Apply Monte Carlo simulations combined with bootstrap methods to evaluate the quality of inference on β_1 using serially correlated data.

Steps

1. Simulate data according to `simulate_regression_with_ar1_errors`.
2. Calculate bootstrap standard errors.
3. Construct a 95% confidence interval for β_1 using both the bootstrap and the theoretical standard errors.
4. Perform Monte Carlo simulations for $T = 100$ and $T = 500$, and assess the empirical coverage of the confidence intervals.